

CS 744: GOOGLE FILE SYSTEM

Shivaram Venkataraman

Fall 2019

ANNOUNCEMENTS


- Assignment 1 out later today
- Group submission form
- Anybody on the waitlist?

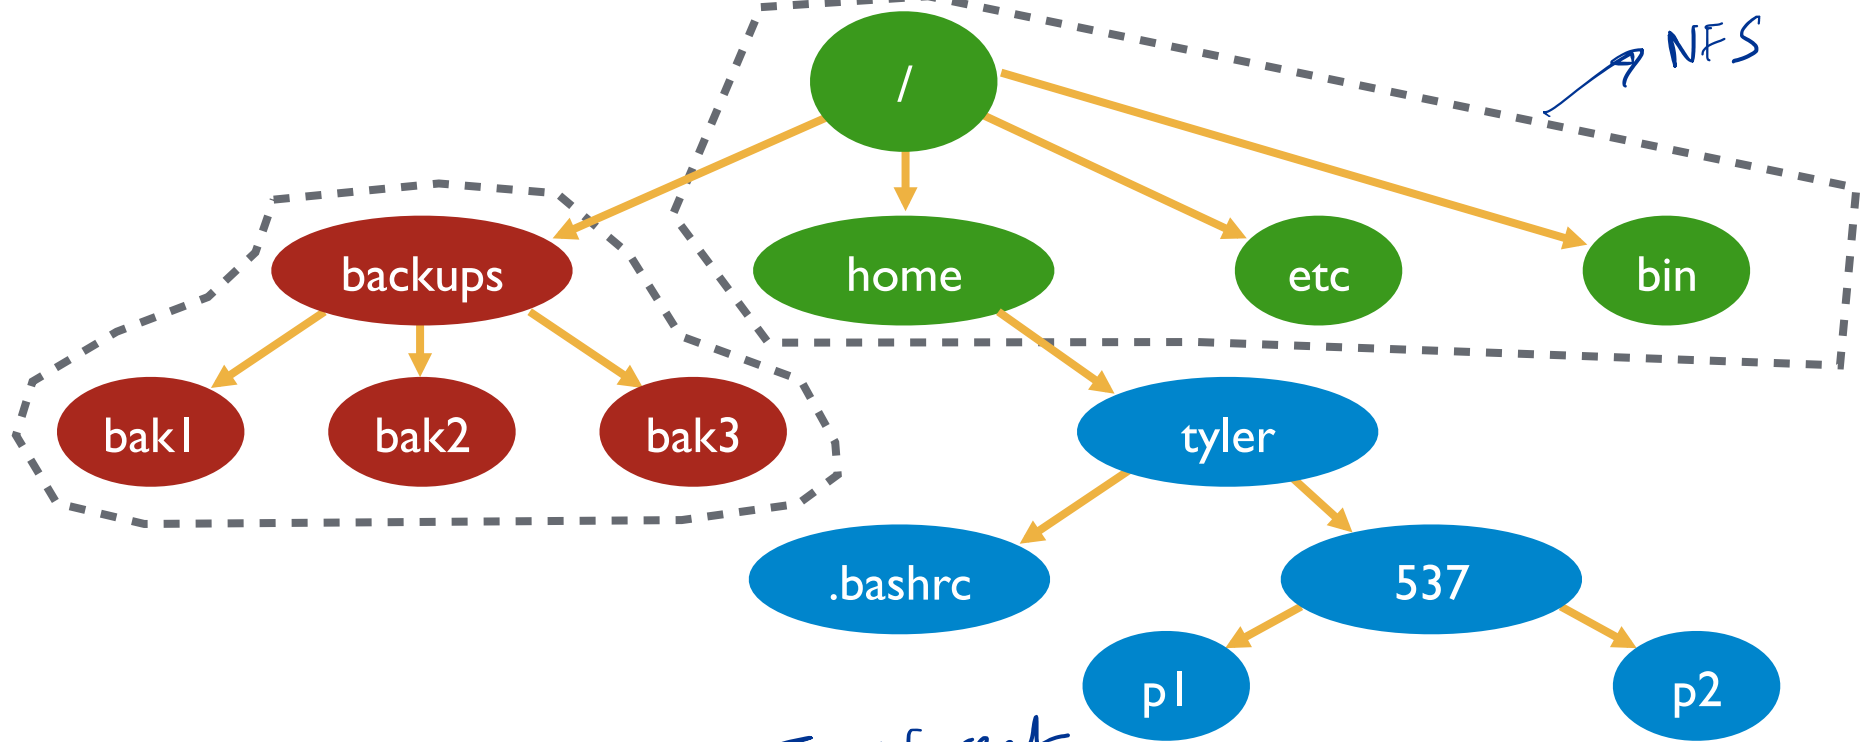
OUTLINE

1. Brief history
2. GFS
3. Discussion
4. What happened next?

HISTORY OF DISTRIBUTED FILE SYSTEMS

SUN NFS



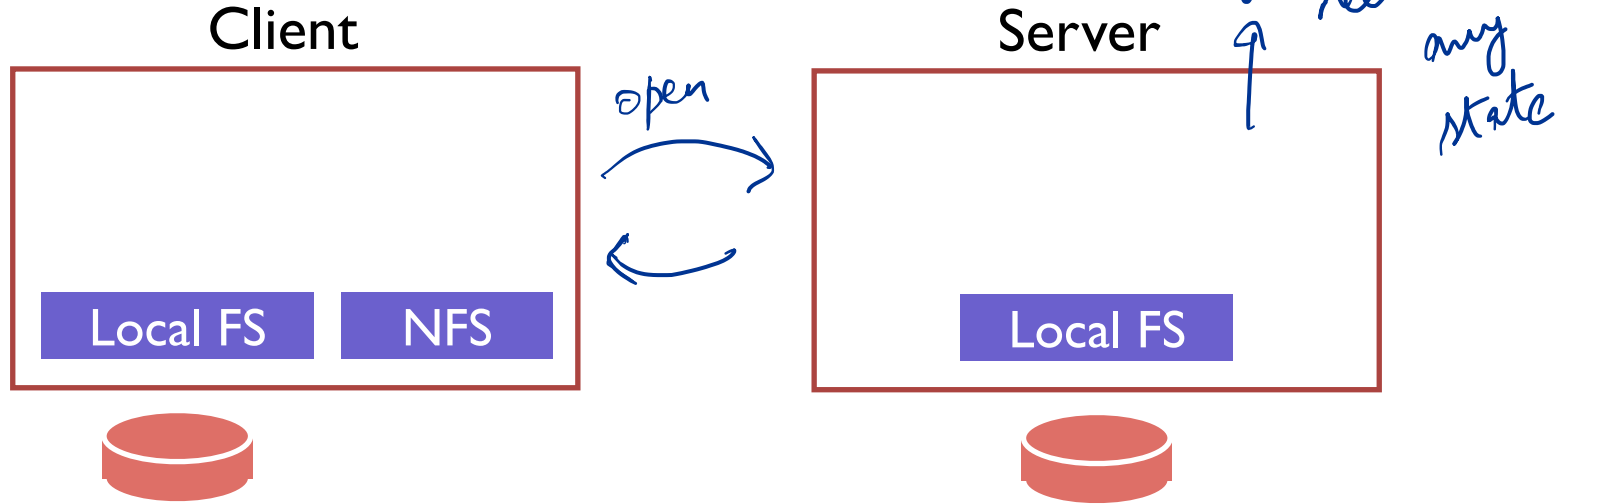


/dev/sda1 **on** /

/dev/sdb1 **on** /backups

NFS **on** /home

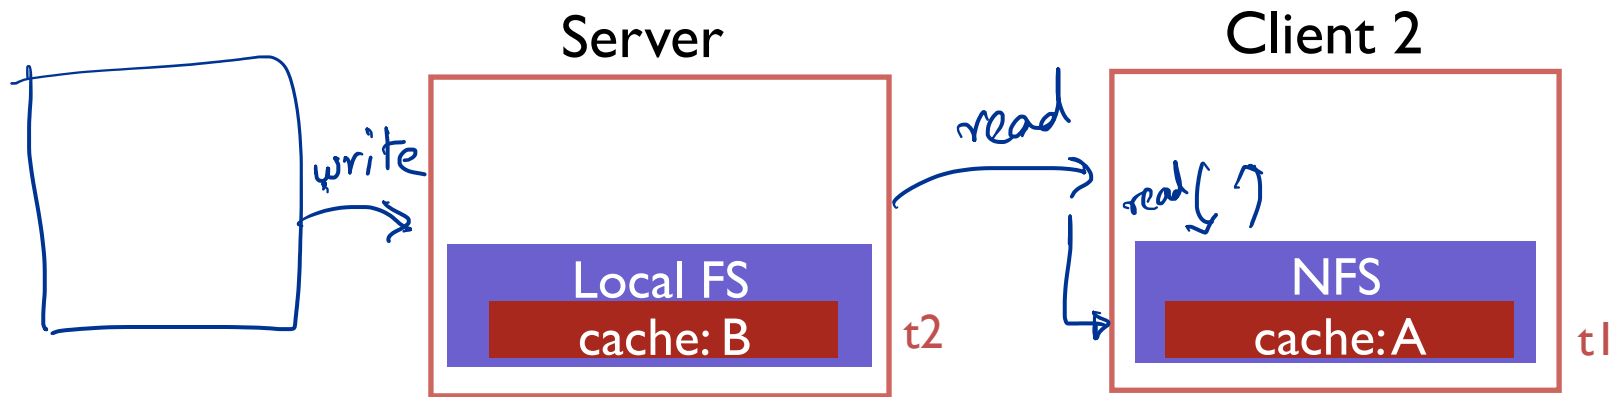
FILE HANDLES



$fd \geq \text{open}("/tmp/foo")$

Examples
open
read

CACHING



Client cache records time when data block was fetched (t_1)

Before using data block, client does a STAT request to server

- get's last modified timestamp for this file (t_2) (not block...)
- compare to cache timestamp
- refetch data block if changed since timestamp ($t_2 > t_1$)

check last updated
timestamp

NFS FEATURES

~ 1985 or
so

NFS handles client and server crashes very well; robust APIs that are:

- stateless: servers don't remember clients
- idempotent: doing things twice never hurts

Caching is hard, especially with crashes

Problems:

- Consistency model is odd (client may not see updates until 3s after file closed)
- Scalability limitations as more clients call stat() on server

ANDREW FILE SYSTEM

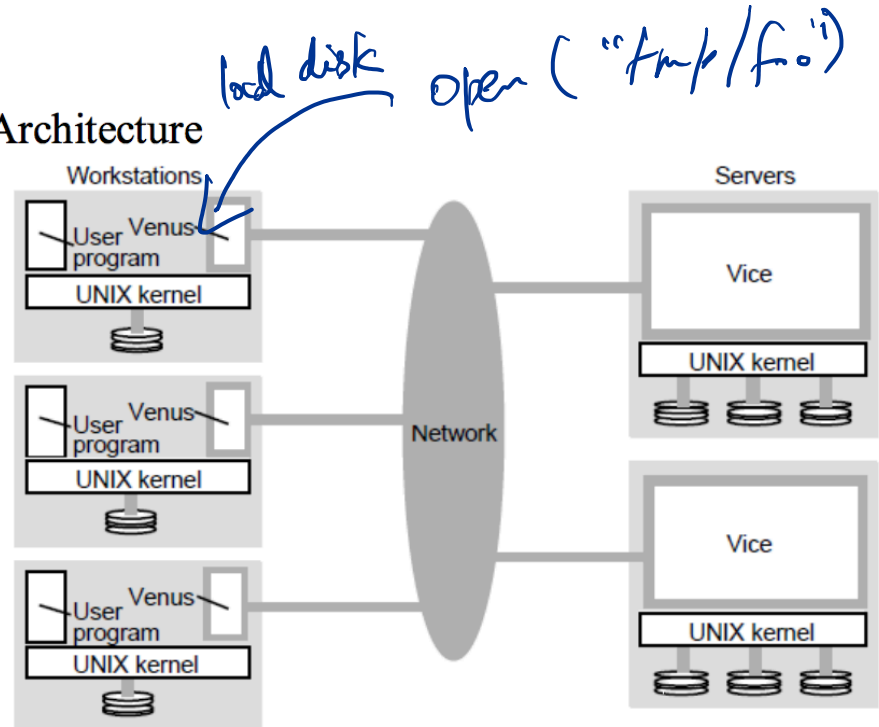
posix API

- Design for scale

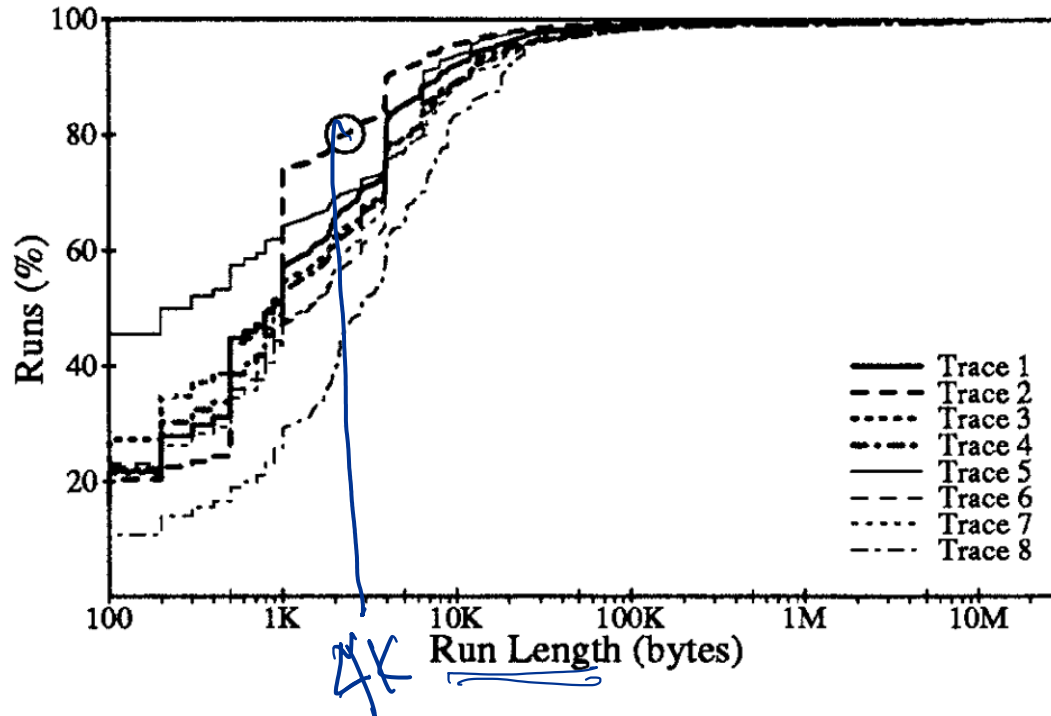
- Whole-file caching

- Callbacks from server

Architecture



WORKLOAD PATTERNS (1991)



Mary G. Baker, John H. Hartman, Michael D. Kupfer, Ken W. Shirriff, and John K. Ousterhout

WORKLOAD PATTERNS (1991)


File Usage	Type of Transfer	Accesses (%)		Bytes (%)	
Read-only	Whole-file	78	(64-91)	89	(46-96)
	Other sequential	19	(7-33)	5	(2-29)
	Random	3	(1-5)	7	(2-37)
Write-only	Whole-file	67	(50-79)	69	(56-76)
	Other sequential	29	(18-47)	19	(4-27)
	Random	4	(2-8)	11	(4-41)
Read/write	Whole-file	0	(0-0)	0	(0-0)
	Other sequential	0	(0-0)	0	(0-0)
	Random	100	(100-100)	100	(100-100)

OCEANSTORE/PAST

Wide area storage systems

Fully decentralized

Built on distributed hash tables (DHT)



Fault tolerance

bench of machines can fail

Co-design
with apps



Workload pattern
changed

GFS: WHY?

- File size > Before
Access size

- Not too many
random writes

→ No caching

↳ Overhead /
Scalability

limits
from
caching?

Components with failures

Files are huge !

GFS: WHY ?

Applications are different

GFS: WORKLOAD ASSUMPTIONS

“Modest” number of large files

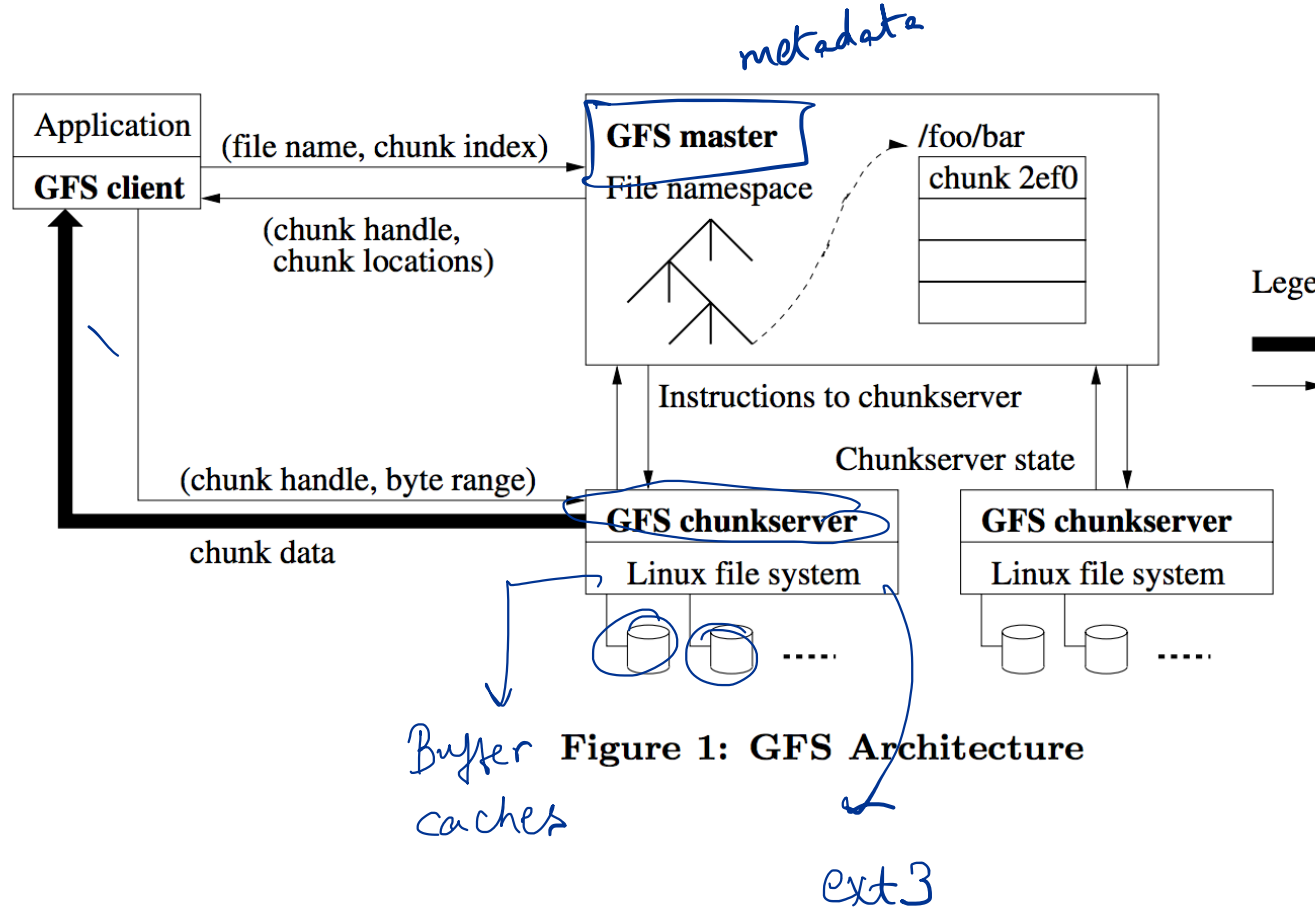
Two kinds of reads: Large Streaming and small random

Writes: Many large, sequential writes. No random

High bandwidth more important than low latency

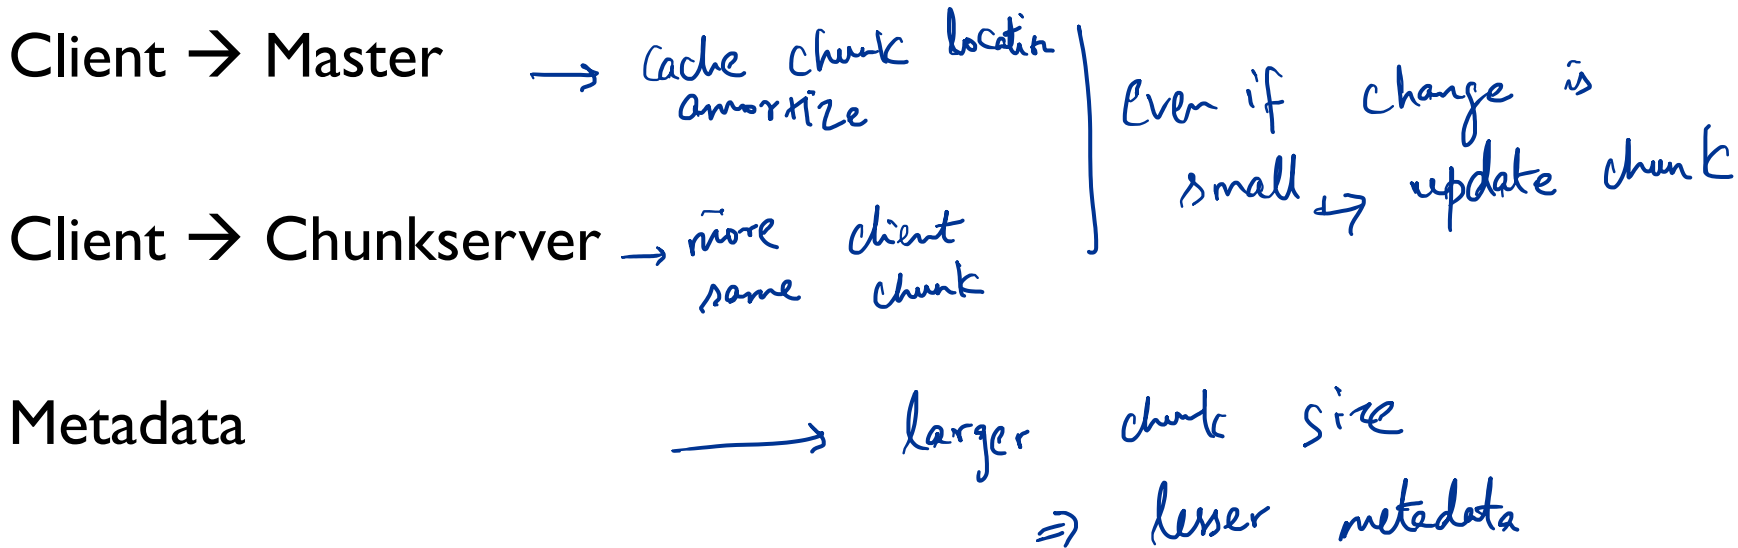
GFS: DESIGN

- Single Master for metadata
- Chunkservers for storing data
- No POSIX API !
- No Caches!

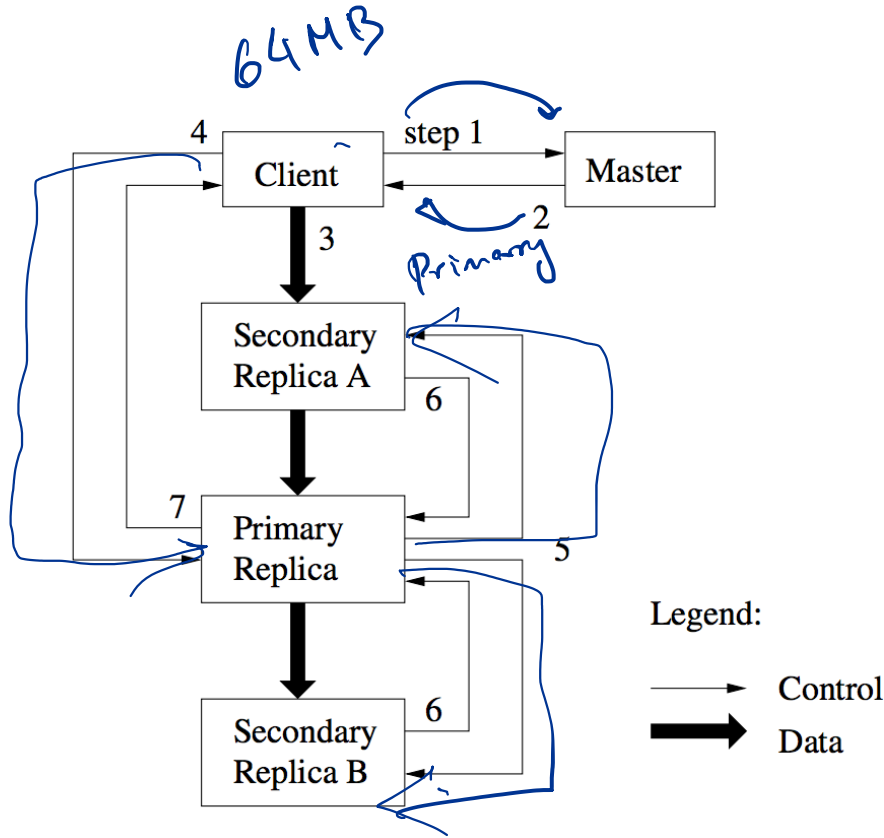



CHUNK SIZE TRADE-OFFS

64 MB



GFS: REPLICATION



- 3-way replication to handle faults
- Primary replica for each chunk
- Chain replication (consistency)
- Dataflow: Pipelining, network-aware

RECORD APPENDS

Write

Client specifies the offset

Record Append

GFS chooses offset



Consistency

At-least once

Atomic

Application level

entire record as a sequence




MASTER OPERATIONS

- No “directory” inode! Simplifies locking
- Replica placement considerations
 - Cross racks for FT
 - Balance writes → balancing empty chunks
 - Disk utilization
- Implementing deletes
 - Lazy deletes using rename
 - Background process
 - recover, manage load, low latency

FAULT TOLERANCE

- Chunk replication with 3 replicas
- Master
 - Replication of log, checkpoint
 - Shadow master
- Data integrity using checksum blocks



Chunk locations → "accurate"


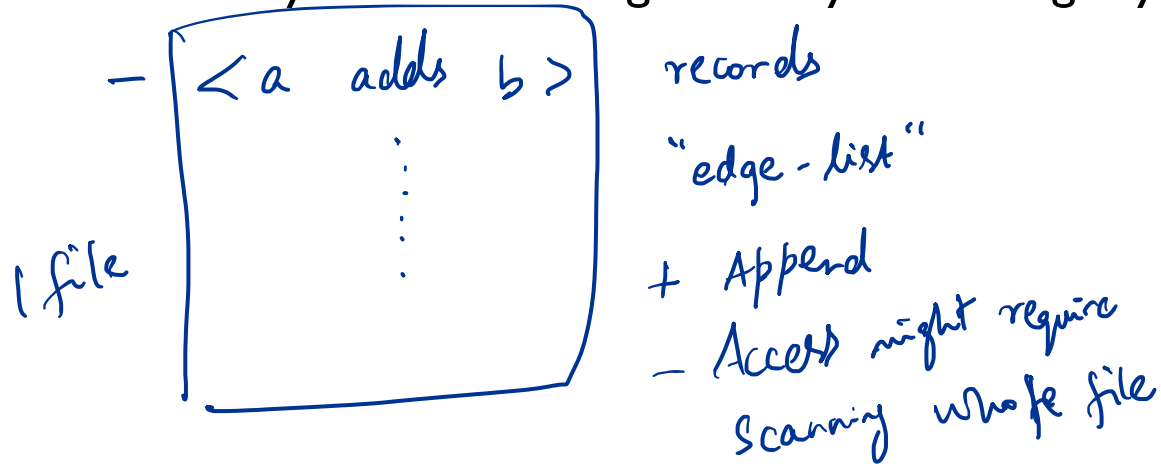
DISCUSSION

GFS SOCIAL NETWORK

You are building a new social networking application. The operations you will need to perform are

- (a) add a new friend id for a given user
- (b) generate a histogram of number of friends per user.

How will you do this using GFS as your storage system ?




100 Mbps \approx 12.5 MB/s
 \approx 76 MB/s Disk

GFS EVAL


List your takeaways from "Figure 3: Aggregate Throughputs"

60-100 MB/s

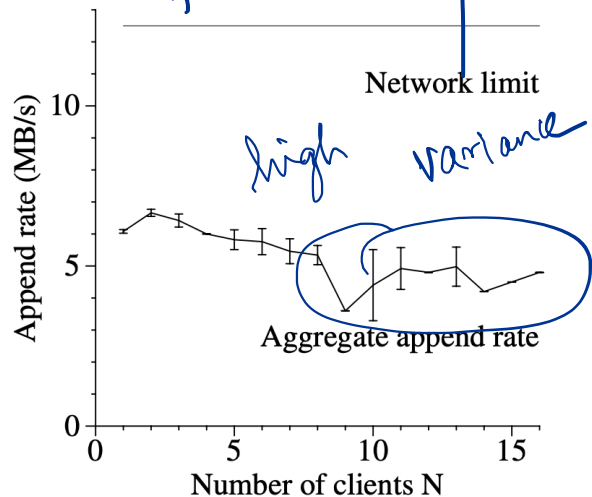
30 N files



(a) Reads



(b) Writes



(c) Record appends

1 file
5-6 MB/s
N clients
last chunk

GFS SCALE

The evaluation (Table 2) shows clusters with up to 180 TB of data. What part of the design would need to change if we instead had 180 PB of data?

64 MB chunks \Rightarrow Metadata explosion!



Chunk size larger

WHAT HAPPENED NEXT



Cluster-Level Storage @ Google

How we use *Colossus* to improve storage efficiency

Denis Serenyi

Senior Staff Software Engineer

dserenyi@google.com

Keynote at PDSW-DISCS 2017: 2nd Joint International Workshop On Parallel Data Storage & Data Intensive Scalable Computing Systems

GFS EVOLUTION

Motivation:

- GFS Master

 - One machine not large enough for large FS

 - Single bottleneck for metadata operations (data path offloaded)

 - Fault tolerant, but not HA

- Lack of predictable performance

 - No guarantees of latency

 - (GFS problems: one slow chunkserver -> slow writes)

GFS EVOLUTION

GFS master replaced by Colossus

Metadata stored in BigTable

Recursive structure ? If Metadata is $\sim 1/10000$ the size of data

100 PB data \rightarrow 10 TB metadata

10TB metadata \rightarrow 1 GB metametadata

1 GB metametadata \rightarrow 100KB meta...


GFS EVOLUTION

Need for Efficient Storage

Rebalance old, cold data

Distributes newly written data evenly
across disk

Manage both SSD and hard disks



HETEROGENEOUS STORAGE



F4: Facebook

Blob stores



Key Value Stores

NEXT STEPS

- Assignment 1 out tonight!
- Next week: MapReduce, Spark