CS 744: RAY

Shivaram Venkataraman

Fall 2019

ADMINISTRIVIA

- Assignment 1 Grades
- Assignment 2 due on Fri
- Course Project emails

Machine Learning

Bismarck

Supervised learning, Unified Interface
Shared memory, Model fits in memory

Parameter Server

Large datasets, large models (PB scale)
Consistency model, Fault tolerance

Tensorflow

Need for flexible programming model
Dataflow graph
Heterogeneous accelerators

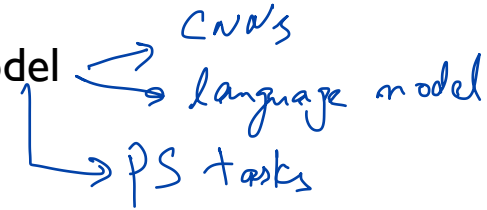

WORKLOADS

Bismarck

→ Convex optimization

→ Small datasets

simple models - SVMs, Logistic reg.

Parameter Server

→ Large datasets → Sparse data

→ High dim parameters → Ad Click model

Tensorflow

↳ Dense features

↳ Advanced models Supervised learning

↳ Image Classification Inception v3

Deep learning

Supervised learning

Training Data labels

Model

fit

REINFORCEMENT LEARNING


002 5 1

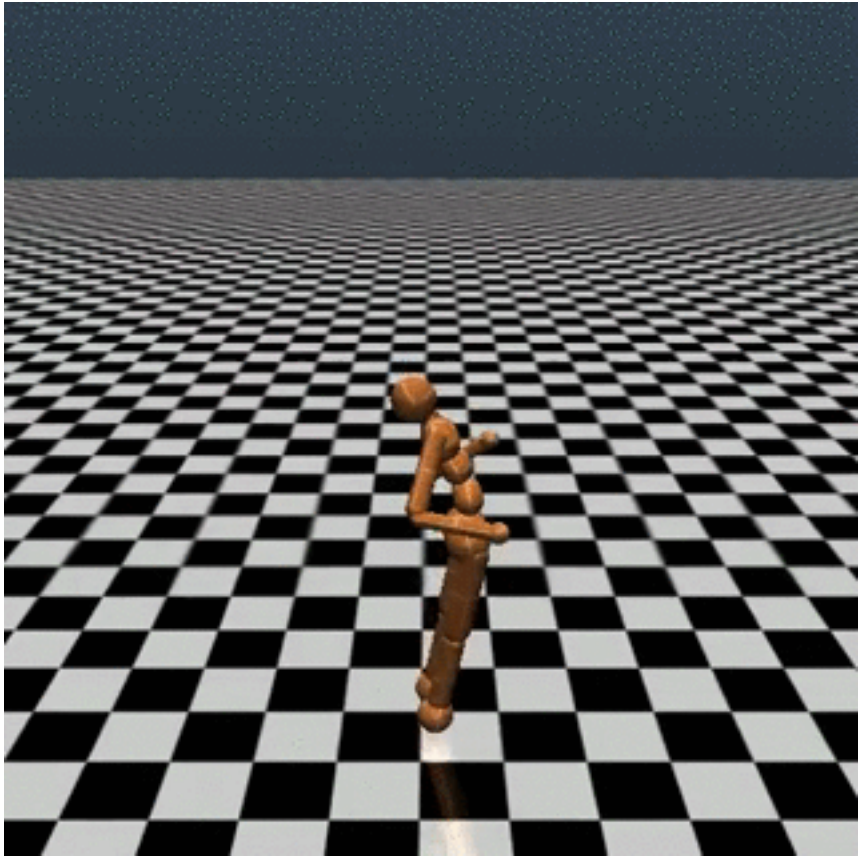

RL SETUP

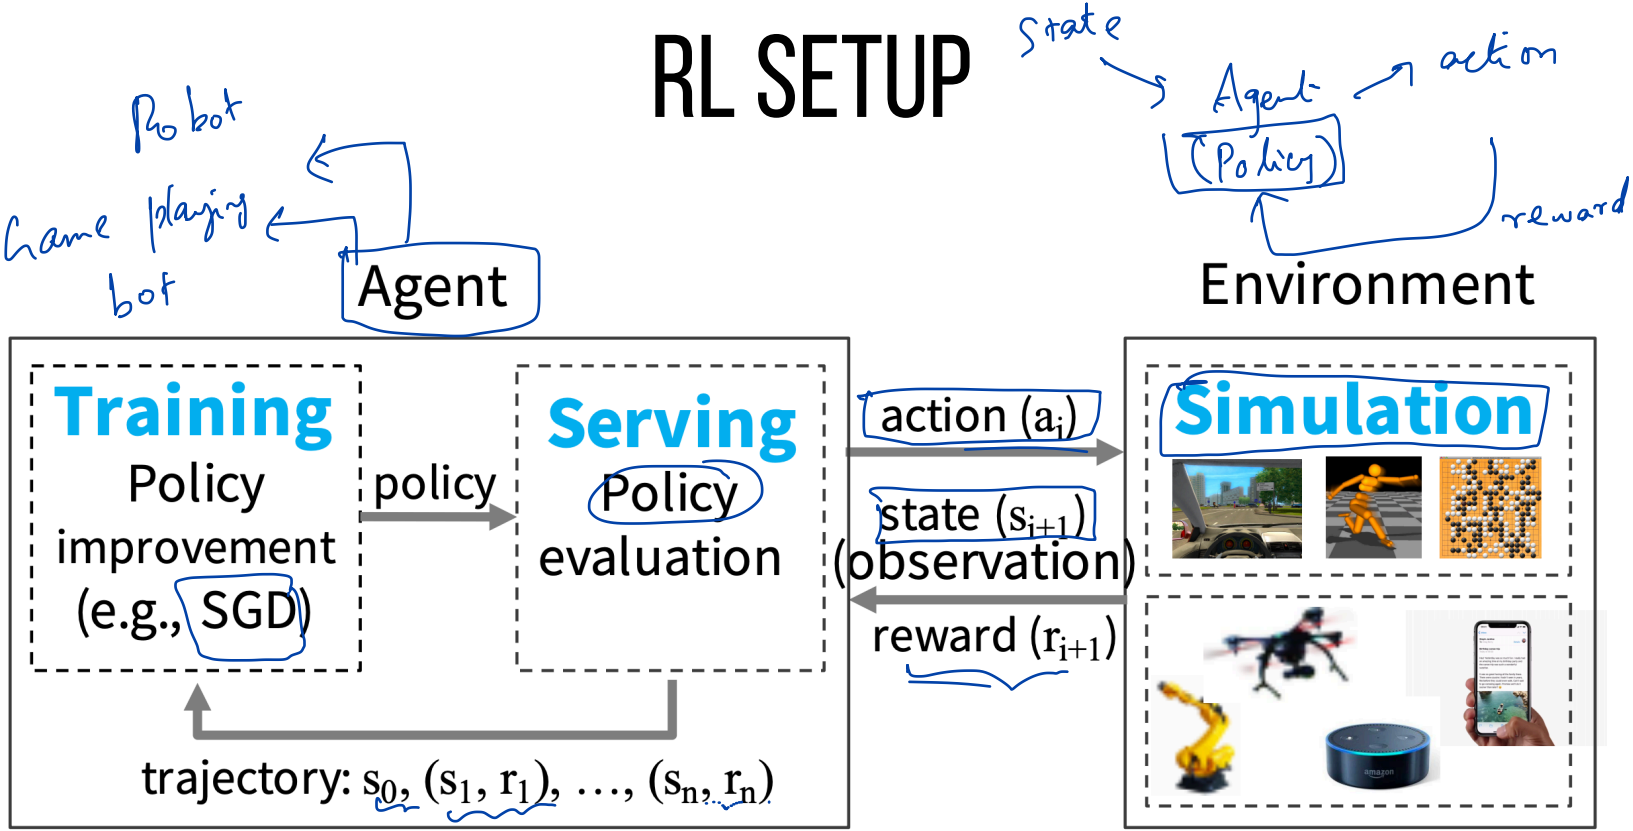

improve Policy given sequence of state, reward

RL REQUIREMENTS

Stateful ≡ Game engine
+
Stateless

Simulation — rollout

↳ Tasks could be of varying length 1ms → seconds / mins

Training — Tasks are deterministic (state, action) → reward

Dynamic execution ≡ iteration structure depends on current iteration

Serving

↳ Very low latency ≡ Parallelism

Rollouts happen in parallel

RAY API

f1 : f.remote (args 1)

Tasks

normal python function

```
futures = f.remote(args)
```

Task which will run f with args

→ Handle to result of task

```
objects = ray.get(futures)
ready = ray.wait(futures, k,timeout)
```

Stateless tasks

Actors

Erlang

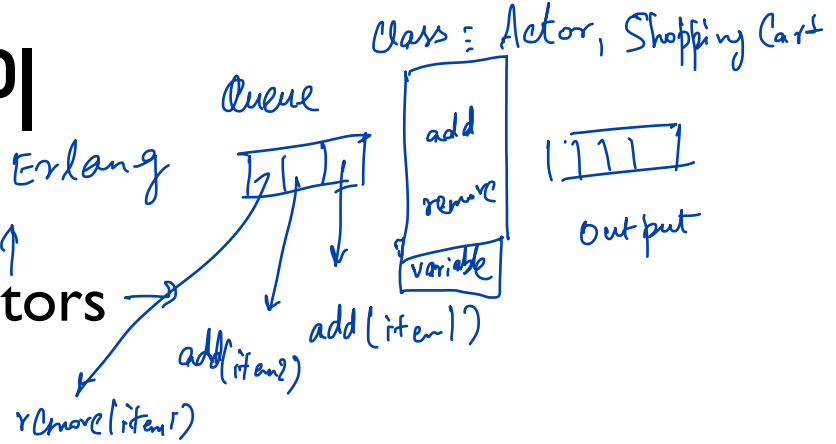


```
actor = Class.remote(args)
futures = actor.method.remote(args)
```

= actor. method. remote (args 1)

f (args) before f (args 1)

COMPUTATION MODEL

Vertices:

Data.

Tasks / Actors

def task()

= task1.remote(a)

Edges

--> Control edge

⟶ Data edge Data vertex Args

Task / Actor

⟹ Stateful edge to ensure ordering

Actor methods

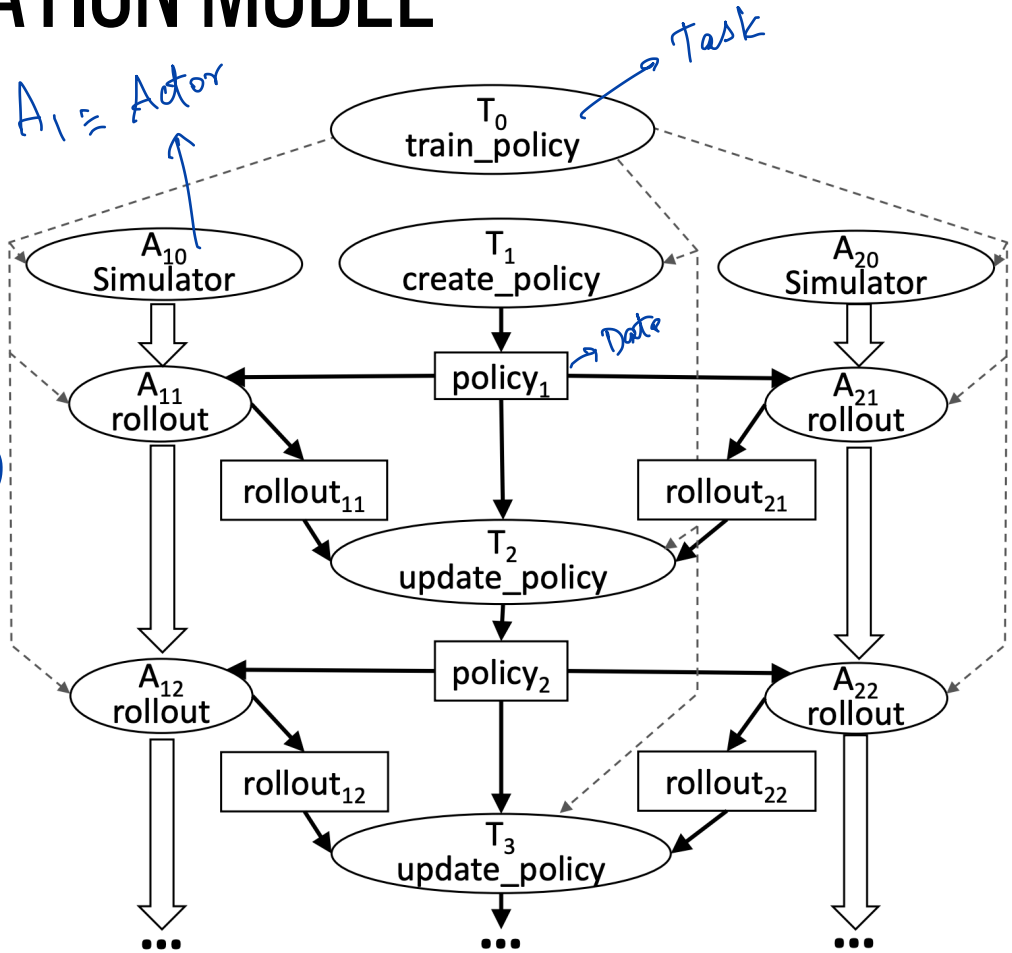

ARCHITECTURE

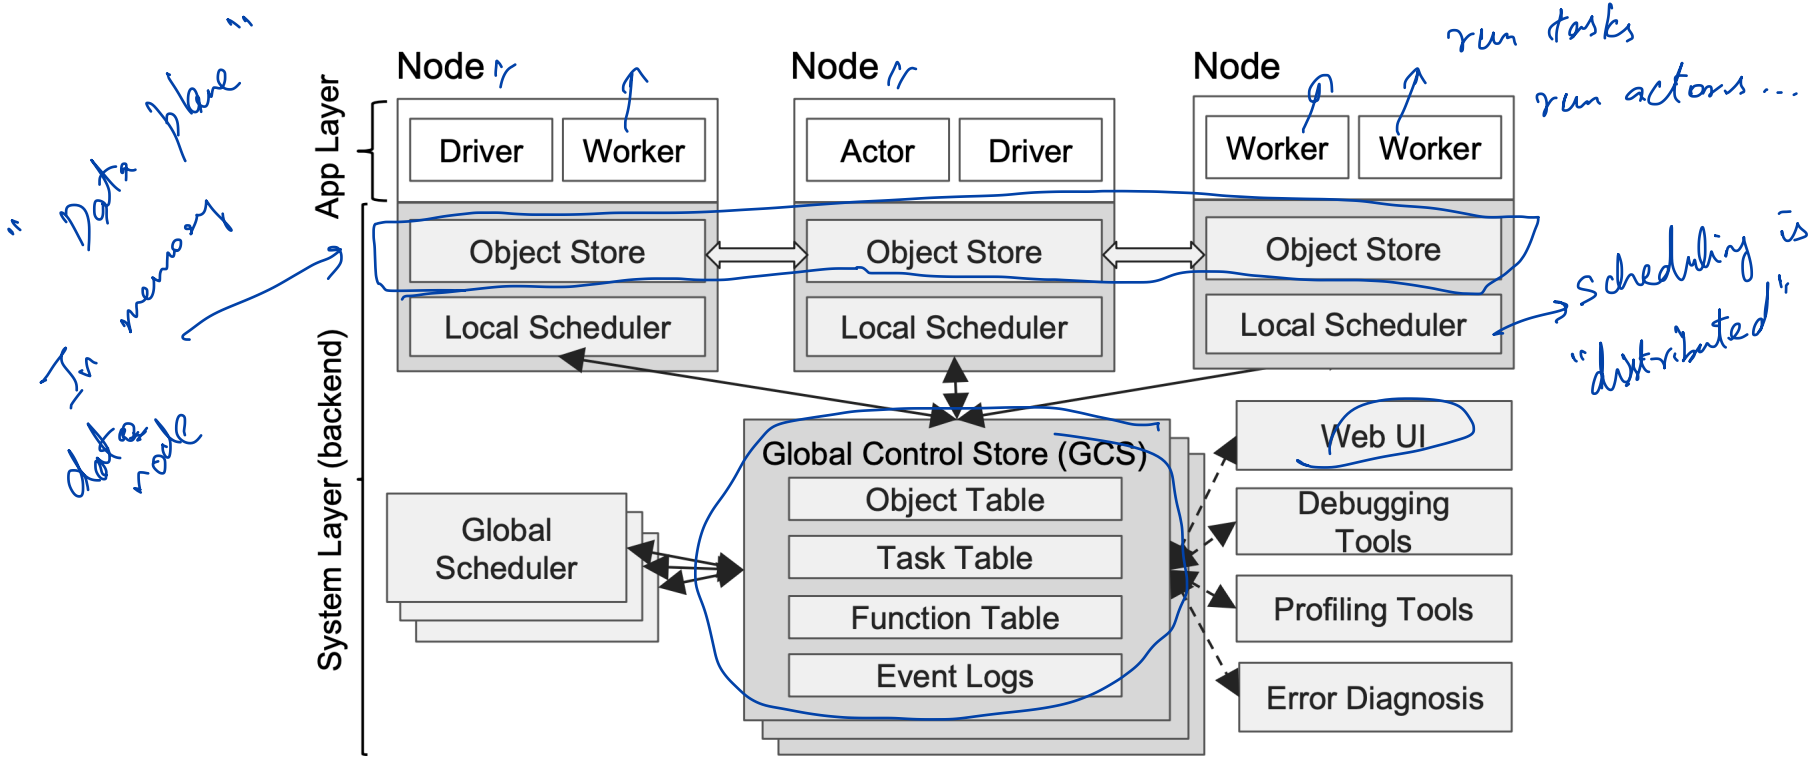

GLOBAL CONTROL STORE

Object table

↳ list of object their locations | Namenode metadata

Task table

↳ lineage, tasks created edge in Comp graph

Function table

↳ Code blocks that
↳ are running

Externalized State

↳ Bottlenecks not in global system

↳ Sharding

Fault tolerance

↳ Chain replication

RAY SCHEDULER

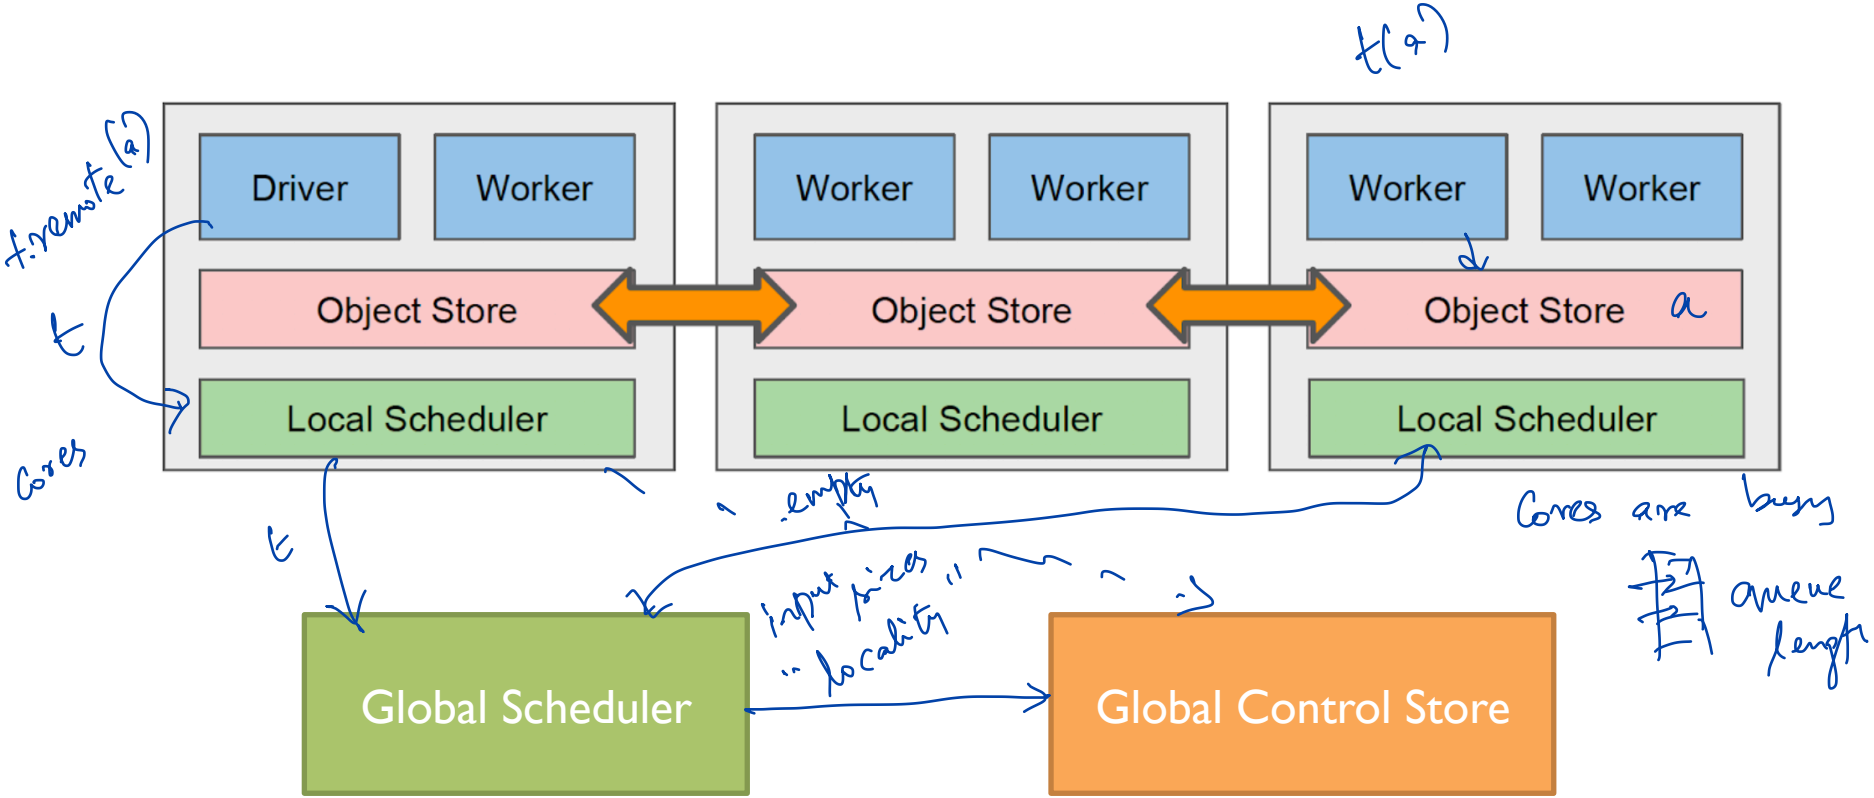

FAULT TOLERANCE

Tasks
↳ lineage from GCS

Actors
↳ "checkpointing"

GCS
↳ sharded, replicated

Scheduler → Stateless,

Driver
↳ computation fails

GCS → extra resources
↳ scalability

Object = Shard

multiple schedulers
or
backup sched.

DISCUSSION

https://forms.gle/QQyLbwjAufJNXWnr6

Consider you are implementing two ~~task~~ apps: a deep learning model training and a sorting application. When will use tasks vs actors and why ?

Actors

wait for other tasks ?

Tasks

You can parallelize and stateless operations

Sorting :

Model training

Parameters are "state"

don't need to broadcast

Flexible synt.

fine grained recovery ?

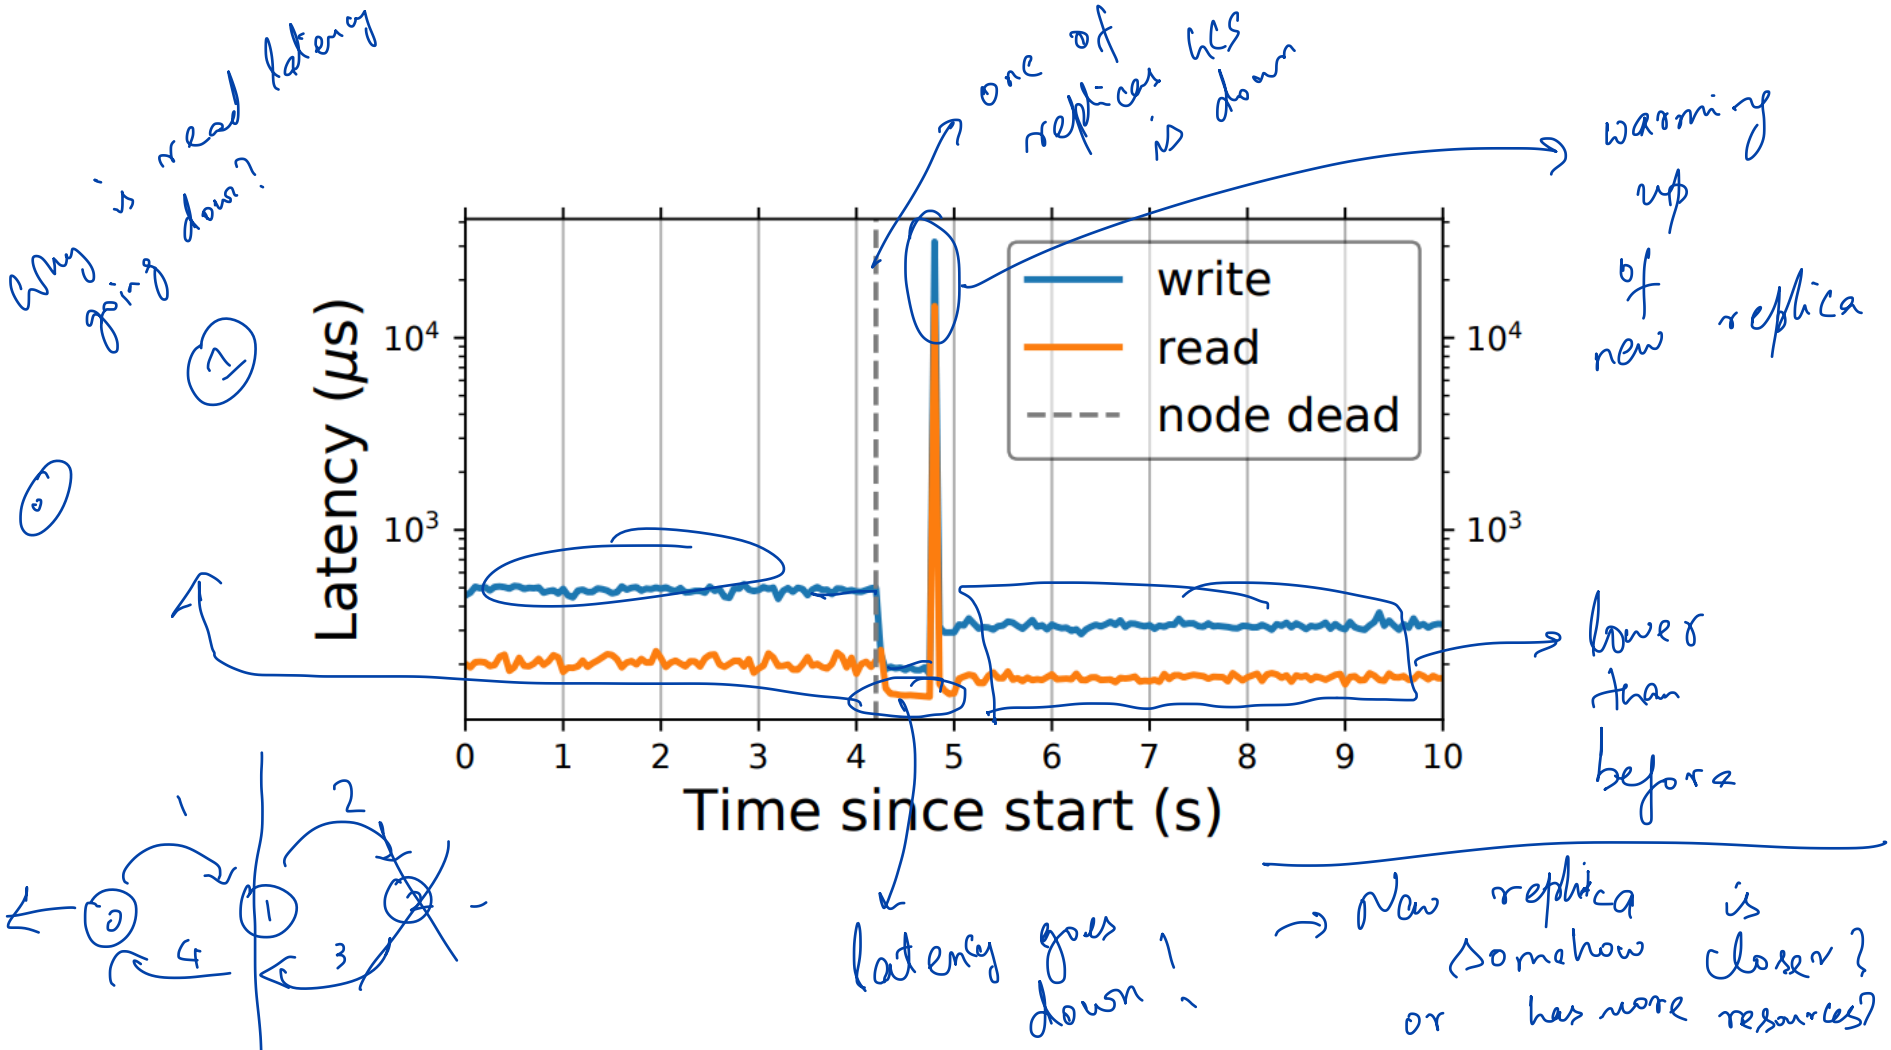
write
read
node dead
Latency (μs)
10^4
10^3
0
1
2
3
4
5
6
7
8
9
10
Time since start (s)

Considering AllReduce using MPI as the baseline parallel programming task. Discuss the improvements made by MapReduce, Spark over MPI and discuss if/how Ray further contributes to the comparison.

NEXT STEPS

Next class: Clipper

Assignment 2 due this week!

Course project