CS 744: CLIPPER

Shivaram Venkataraman

Fall 2019

ADMINISTRIVIA

- Assignment 2 grading
- Midterm details
- Course Project template - Due Thu

MACHINE LEARNING SO FAR

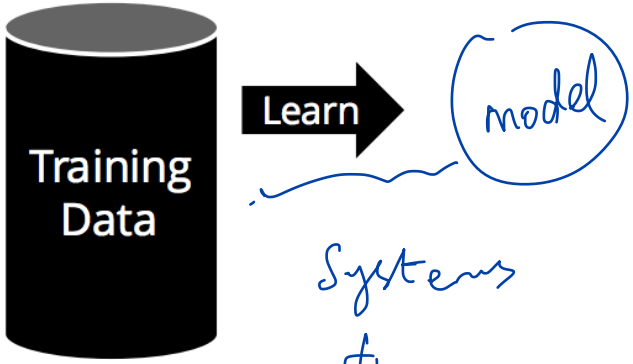

Systems to learn models

Shared memory

Distributed

Flexible Programming Model

Reinforcement learning

MACHINE LEARNING: INFERENCE

Movie Recommendation

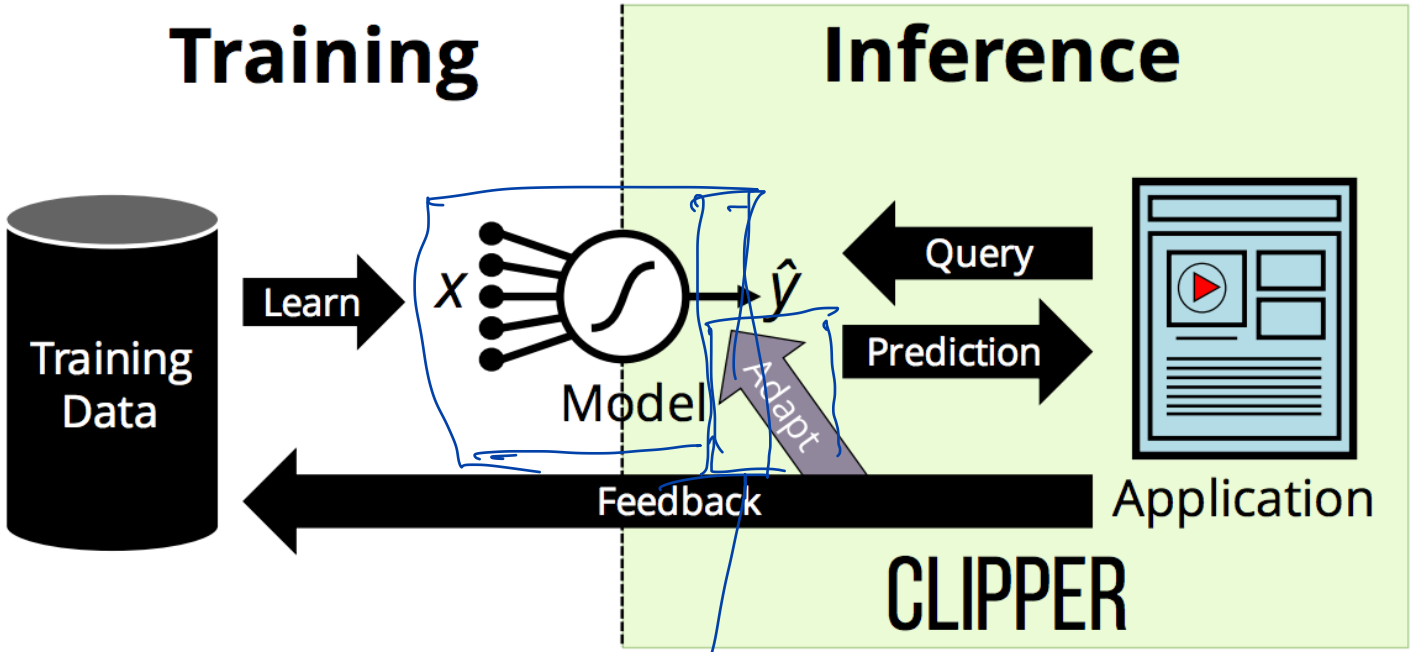

Incrementally update the model/Prediction

GOALS

- Interactive latencies (tail latency < 100ms) → page load times
- High throughput to handle load → % handle many users

  → general
- Improved prediction accuracy
- Generality (?) → support multiple ML frameworks

ARCHITECHTURE

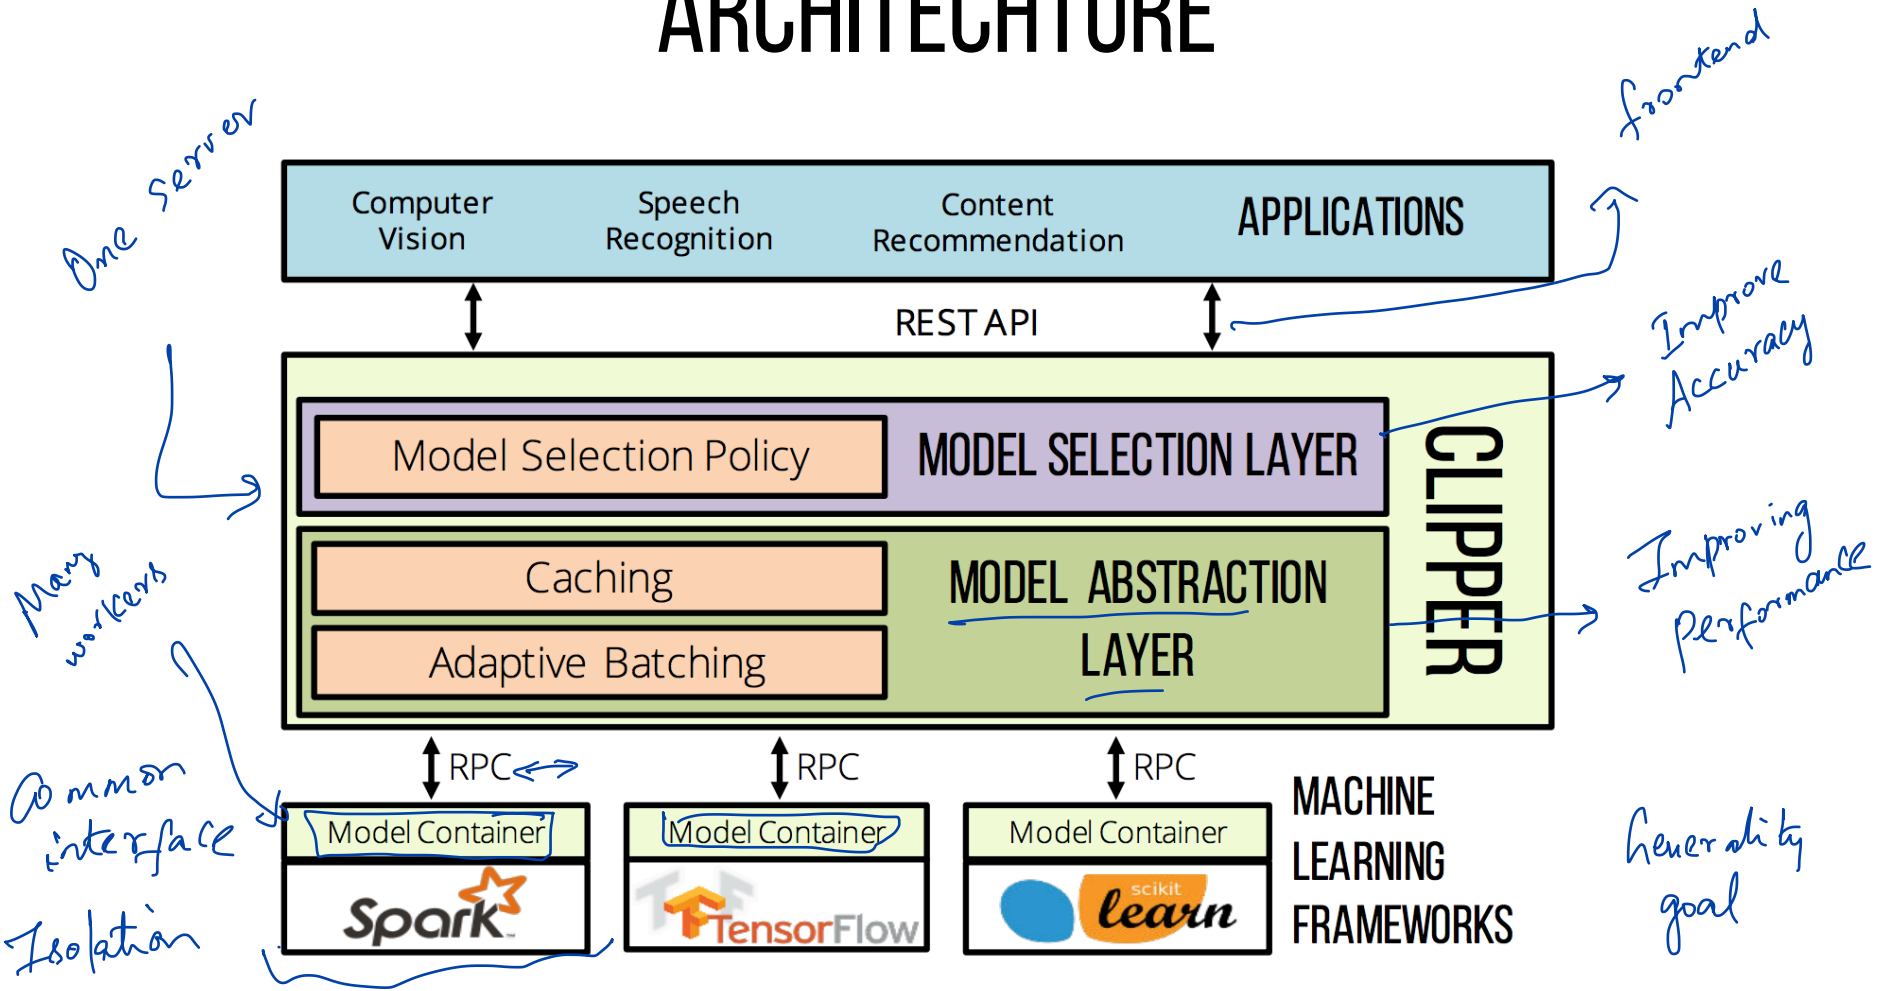

MODEL CONTAINERS

→ No workers / PS tasks inside MC

→ No training data

→ Assumption: weights fit in MC

```
interface Predictor<X,Y> {
  List<List<Y>> pred_batch(List<X> inputs);
}
```

for each input

list of outputs

Training → labels

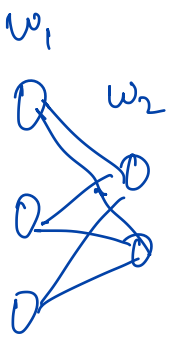

- Run using Docker containers

- Can be replicated across machines

↳ Replication factor could vary

Pre-processing + Model inference

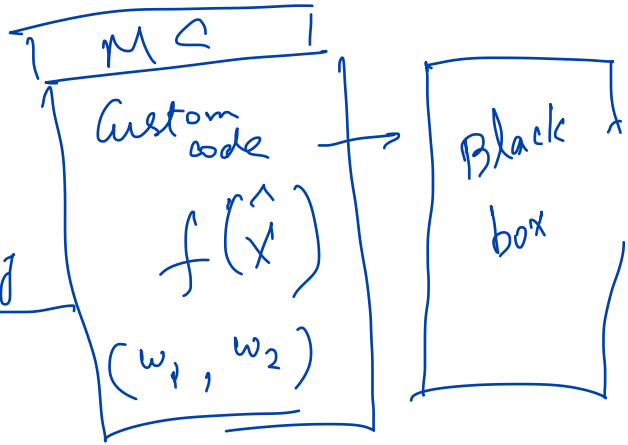

MODEL ABSTRACTION LAYER

Caching

- Improve performance for frequent queries
- LRU eviction policy
- Important for feedback

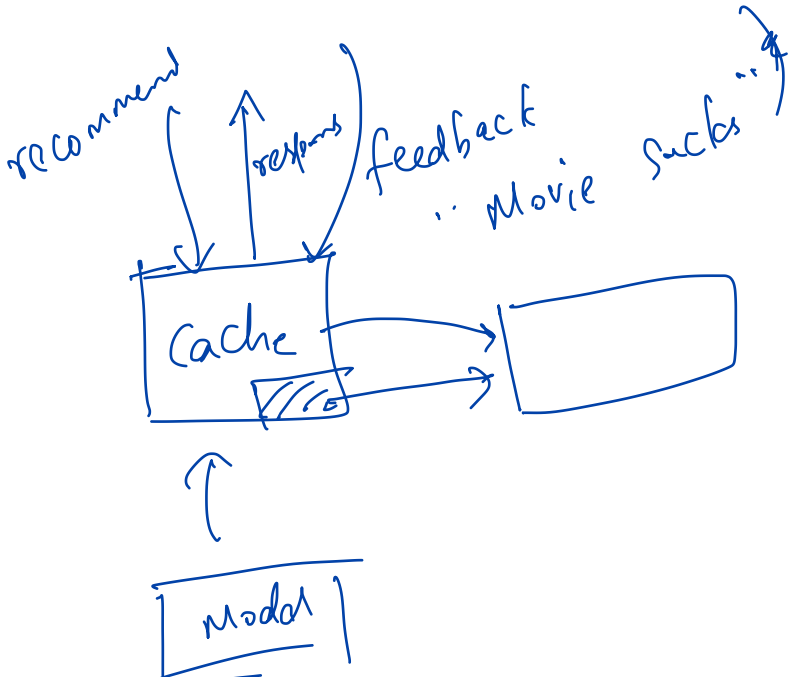

100 request
batch size : 50
= 100 RPCs w/o batching 2 RPCs w/ batching

BATCHING, QUEUING

Goals, Insight

- Increase latency (within SLO) for improved throughput
- Reduce RPC overheads
- GPU / BLAS acceleration

→ Opt Batch size different for each container

→ 500 batchsize : 20 ms

Approach

- Per container queues.
- Maximum batch size.
- Why?

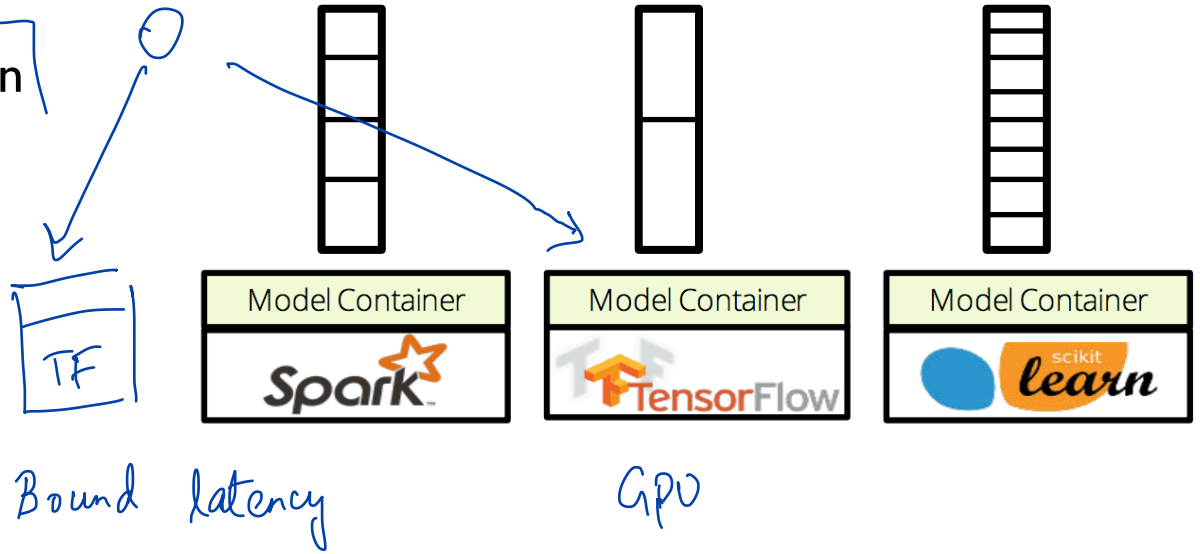

→ Bound latency

GPU

ADAPTIVE BATCHING

AIMD: Additive Inc Multiplicative Dec

Why ?

Runtime variations

Delayed: Wait until batch exists

Why?

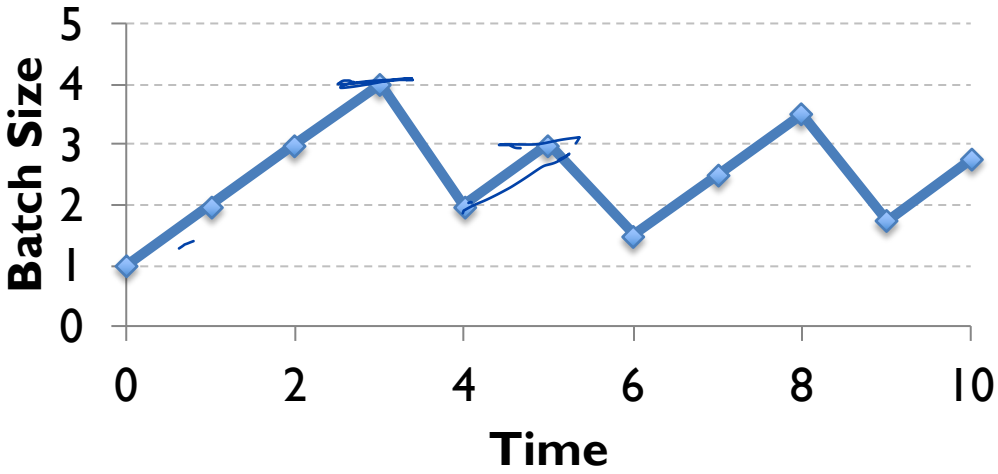

MODEL SELECTION

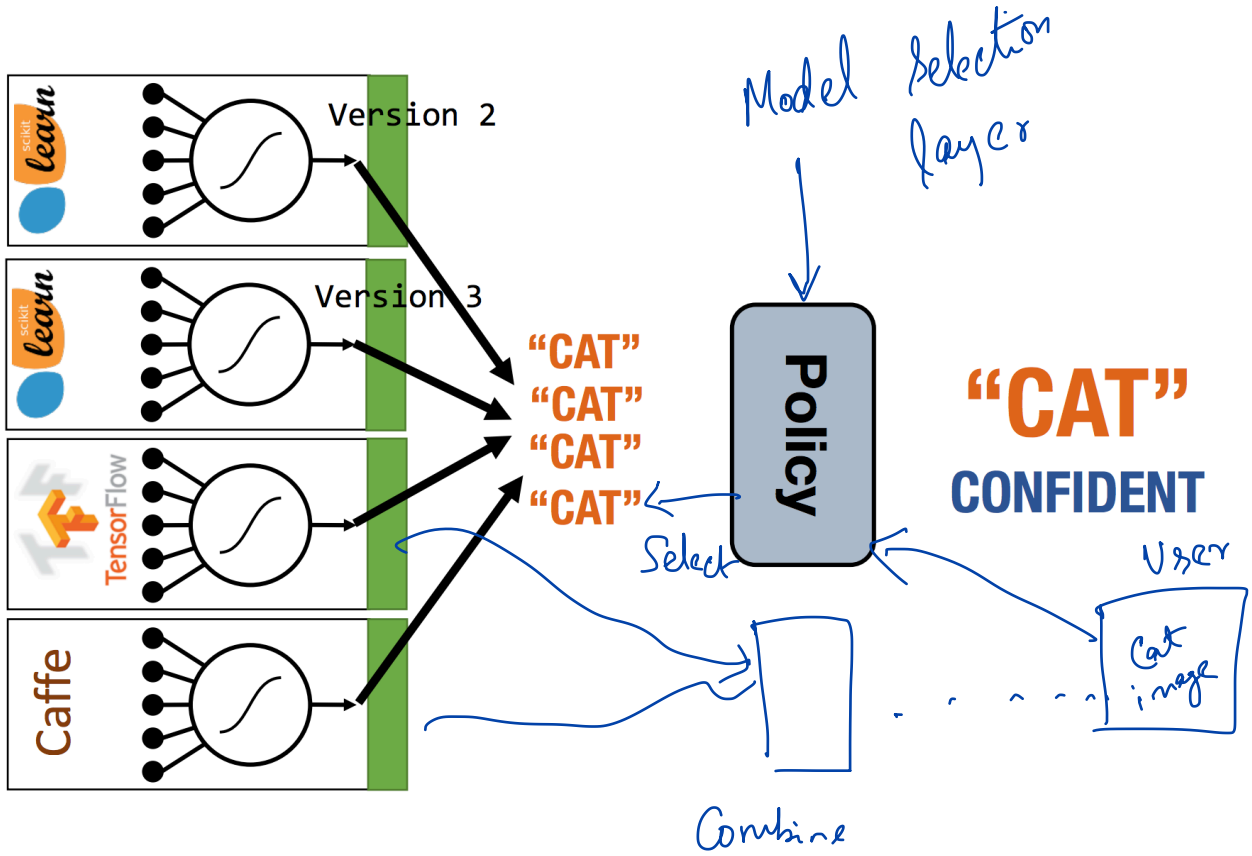

SINGLE MODEL SELECTION

Multi-Arm Bandit formulation

- Explore vs Exploit
- Regret: Loss by not picking optimal action
- Goal: Minimize regret

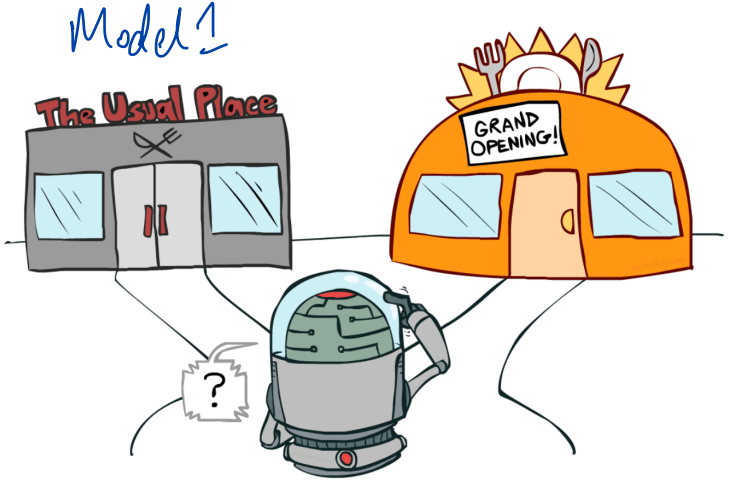

Clipper

- Exp3 algorithm → uses feedback to influence next selection.
- Single evaluation
- Scales to more models

MULTI MODELS

Ensemble

- Combine output from models (weighted average)
- How do we get the weights ? → linear Combination

Robust Prediction

- React to model changes
- Output confidence score

STRAGGLER MITIGATION

Why do stragglers occur?

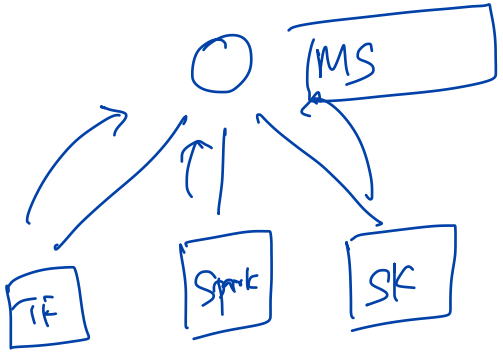

All of the requests finish

Approach

↳ Approx result not exact

↳ Better approx than late!

ML specific

Slower model could be more accurate!

Deadline → Pick the best based on results within deadline

TAKEAWAYS

- ML inference: Workloads + Requirements
- Layered architecture provides generality
- Caching, Batching, Replication to improve latency, throughput
- Multi-Arm bandits to improve accuracy

DISCUSSION

https://forms.gle/pZMuhCWcap2q3LQJ9

(Discussion question from last week)

Considering AllReduce using MPI as the baseline parallel programming task. Discuss the improvements made by MapReduce, Spark over MPI and discuss if/how Ray further contributes to the comparison.

- Ease of programming → MR
  - ↳ Spark
- Fault tolerance → MR → disk
  - Spark → Lineage
  - Ray ↗
- Straggler mitigation → MR → Speculative tasks
- Network Usage → Controlled partition
  - Locality → Mem, disk → MR
    - ↳ Spark, Ray

Scheduling benefits
- ↳ Locality
- ↳ Scale
  - ↳ Ray

Consider a scenario where you run a model serving service that hosts a number of different ~~models~~. The traffic for some ~~models~~ is sporadic (e.g. only a few hours where they are used). What are some advantages / disadvantages of using Clipper for such a service?

application

apps

→ Automatically scale up, scale down

At least 1 container active

→ Batching is not effective, sporadic

→ SLOs could be different →

Caching might not be effective

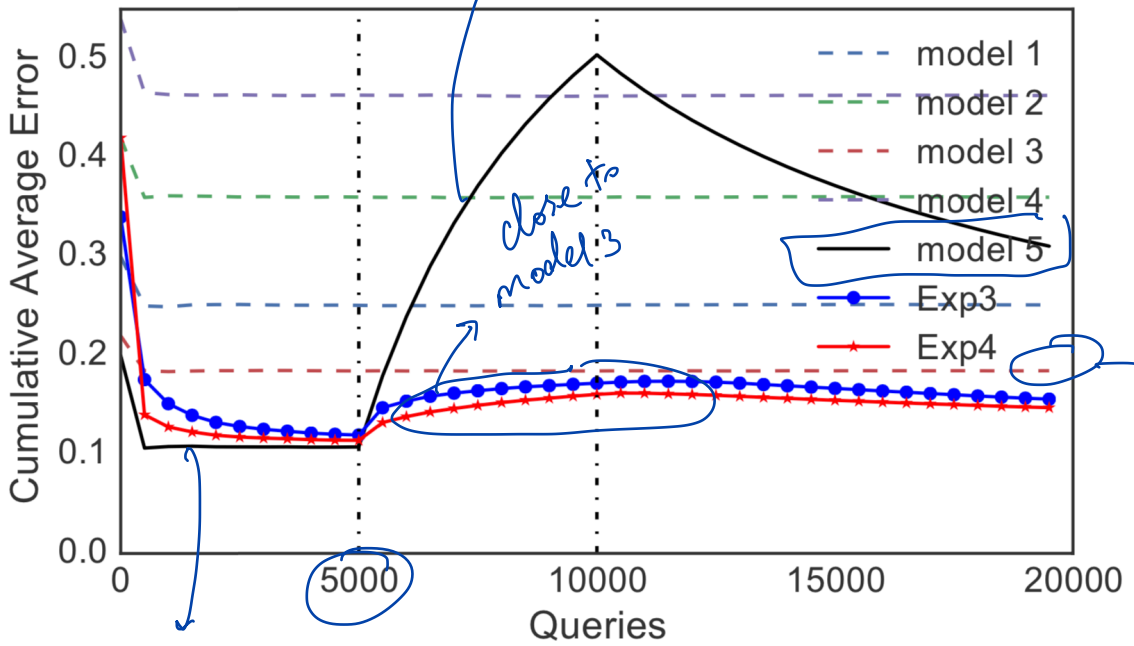

model 5 is bad here

close to model 3

model 1 + 2 + 4 + model 5

model 3

exp3, exp4 model 5 is close to